Assignment 2. Sentiment Analysis

Hwanjo Yu
CSED342 - Artificial Intelligence

Contact: TA Chunghyun Park (p0125ch@postech.ac.kr)
Deadline: 18 Mar 2024 at 2:00 pm. (50% penalty for every 1 day)

General Instructions

You should write answers in `submission.py` between

```
# BEGIN_YOUR_ANSWER
```

and

```
# END_YOUR_ANSWER
```

In **Problem 1** you just need to provide your answers.

In **Problem 2** and **Problem 3** you have to implement some functions. You can add other helper functions outside the answer block if you want. Do not make changes to files other than `submission.py`.

Please use **Python 3.7** and **NumPy** to develop your code. You can refer to the official documentation of Numpy[1] to install it. We recommend you use Conda to set up the environment. You have to write answers in the same format as provided in `submission.py`.

Your code will be evaluated on two types of test cases, **basic** and **hidden**, which you can see in `grader.py`. Basic tests, which are fully provided to you, do not stress your code with large inputs or tricky corner cases. Hidden tests are more complex and do stress your code. The inputs of hidden tests are provided in `grader.py`, but the correct outputs are not. To run all the tests, type

```
python grader.py
```

This will tell you only whether you passed the basic tests. On the hidden tests, the script will alert you if your code takes too long or crashes, but does not say whether you got the

---

[1]https://numpy.org/install

correct output. You can also run a single test (e.g., `3a-0-basic`) by typing

```
python grader.py 3a-0-basic
```

We strongly encourage you to read and understand the test cases, create your own test cases, and not just blindly run `grader.py`.
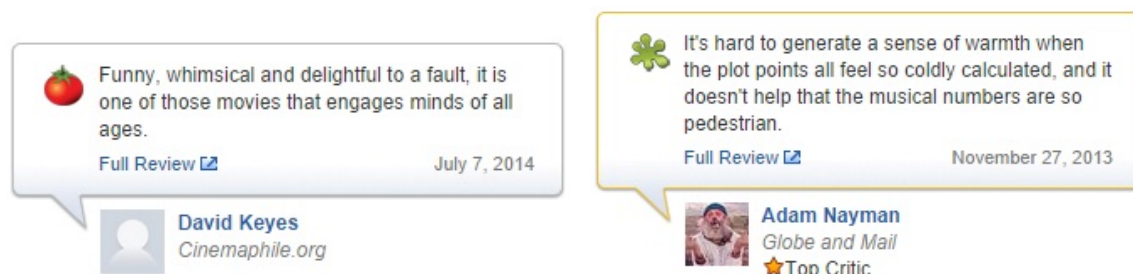

Advice for this homework:

- Words are simply strings separated by whitespace. Don't normalize the capitalization of words (treat *great* and *Great* as different words).

- You might find some useful functions in `util.py`. Have a look around in there before you start coding.

Problems

Problem 1. Hinge Loss

Here are two reviews of "Frozen," courtesy of Rotten Tomatoes (no spoilers!):

Rotten Tomatoes has classified these reviews as "positive" and "negative," respectively, as indicated by the in-tact tomato on the left and the splattered tomato on the right. In this assignment, you will create a simple text classification system that can perform this task automatically.

Problem 1a [2 points]

We'll warm up with the following set of four mini-reviews, each labeled positive ($+1$) or negative (-1):

- ($+1$) so interesting

- ($+1$) great plot

- (-1) so bored

- (-1) not interesting

2

Each review x is mapped onto a feature vector $\phi(x)$, which maps each word to the number of occurrences of that word in the review. For example, the first review maps to the (sparse) feature vector $\phi(x) = \{\text{so} : 1, \text{interesting} : 1\}$. Recall the definition of the hinge loss:

$$\text{Loss}_{\text{hinge}}(x, y, \mathbf{w}) = \max\{0, 1 - \mathbf{w} \cdot \phi(x)y\},$$

where y is the correct label.

Suppose we run stochastic gradient descent, updating the weights according to

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} \text{Loss}_{\text{hinge}}(x, y, \mathbf{w}),$$

once for each of the four examples in order. After the classifier is trained on the given four data points, <u>what are the weights of the six words</u> ('so', 'interesting', 'great', 'plot', 'bored', 'not') that appear in the above reviews? Use $\eta = 1$ as the step size and initialize $\mathbf{w} = [0, ..., 0]$. Assume that $\nabla_{\mathbf{w}} \text{Loss}_{\text{hinge}}(x, y, \mathbf{w}) = 0$ when the margin is exactly 1.

Problem 2: Sentiment Classification

In this problem, we will build a binary linear classifier that reads movie reviews and guesses whether they are "positive" or "negative."

Problem 2a [2 points]

Implement the function *extractWordFeatures*, which takes a review (string) as input and returns a feature vector $\phi(x)$ (you should represent the vector $\phi(x)$ as a *dict* in Python).

Problem 2b [8 points]

We're going to train a linear predictor, which can be represented by a logistic regression model. Here is the definition of linear predict:

$$f_w(x) = \begin{cases} +1 & \text{if } \mathbf{w} \cdot \phi(x) > 0 \\ -1 & \text{if } \mathbf{w} \cdot \phi(x) < 0 \end{cases}$$

$$= \begin{cases} +1 & \text{if } \sigma(\mathbf{w} \cdot \phi(x)) > 0.5 \\ -1 & \text{if } \sigma(\mathbf{w} \cdot \phi(x)) < 0.5 \end{cases}$$

where σ is a logistic(or sigmoid) function.

Your task is to implement the function *learnPredictor* using stochastic gradient descent, minimizing the negative log-likelihood loss (NLL) defined as:

$$\text{Loss}_{\text{NLL}}(x, y, \mathbf{w}) = -\log(p_{\mathbf{w}}(y \mid x))$$

$$p_{\mathbf{w}}(y \mid x) = \begin{cases} \sigma(\mathbf{w} \cdot \phi(x)) & \text{if } y = 1 \\ 1 - \sigma(\mathbf{w} \cdot \phi(x)) & \text{if } y = -1 \end{cases}$$

You should first derive $\nabla_{\mathbf{w}} \text{Loss}_{\text{NLL}}(x, y, \mathbf{w})$, then exploit the formula to update weights for each example. Initialize the weights $\mathbf{w}$ as $\mathbf{0}$. Also, you can print the training error and test error after each iteration through the data, so it's easy to see if your code is working.

Problem 2c [3 points]

The previous features include unigram(single) words only, which cannot consider the context of a word in an utterance. In this task, we'll incorporate n-gram words into features. In other words, the feature vector will count the number of occurrences of consecutive words of length n. Implement *extractNgramFeatures* which extracts n-gram word features.

Problem 3: Multi-Layer Perceptron & Backpropagation

In this problem, we will build a binary non-linear classifier with Multi-Layer Perceptron (MLP) and train the classifier with the NLL loss using stochastic gradient descent. Note that you are not allowed to use Python libraries that support automatic differentiation (*e.g.*, PyTorch, TensorFlow). Use **Numpy** to implement your codes.

Suppose we have a feature extractor ϕ that produces 2-dimensional feature vectors, and a training dataset $\mathcal{D}_{\text{train}} = \{(x_1, y_1), (x_2, y_2), (x_3, y_3), (x_4, y_4)\}$ (a.k.a XOR problem) with

1. $\phi(x_1) = [0, 0]$, $y_1 = 0$

2. $\phi(x_2) = [0, 1]$, $y_2 = 1$

3. $\phi(x_3) = [1, 0]$, $y_3 = 1$

4. $\phi(x_4) = [1, 1]$, $y_4 = 0$.

Problem 3a [4 points]

Implement the *forward* function in *MLPBinaryClassifier* to predict the likelihood of $y = 1$ using 2-dimensional features with sigmoid activation. Do not modify the *init* function of *MLPBinaryClassifier* provided. Note that you need to save intermediate activations within the *forward* function for backpropagation (Problem 3b).

Problem 3b [8 points]

Implement the *loss* and *backward* functions in *MLPBinaryClassifier* to backpropagate the negative log-likelihood (NLL) loss to network parameters: $W1$, $b1$, $W2$, and $b2$. Utilize the saved activations from the forward process to compute the gradient for each parameter.

Problem 3c [2 points]

Implement the *update* and *train* functions in *MLPBinaryClassifier* to train the network using stochastic gradient descent. Within the *train* function, you may need to call the *forward*, *loss*, *backward*, and *update* functions you implemented in the previous problems. Use the batch size of 1 for stochastic gradient descent and do not shuffle the training data within the *train* function.